



File Edit View Insert Cell Kernel Widgets Help
Not Trusted Python 3

Regression Analyses with Cricket Data

In week 1, we took a brief look at the cricket match of statistics of the Indian Premier league in 2018 (IPL2018teams dataset). In this week, we will look at the player level statistics. In particular, we are interested in whether the player performance impact their salaries.

Import useful libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.formula.api as sm
```

Import cricket data

In our data repository, there is a data set "IPL18Player.csv" which contains performance statistics as well as salary information of cricket players in the Indian Premier League in 2018.

```
In [2]: IPLPlayer=pd.read_csv("../Data/Week 4/IPL18Player.csv")
IPLPlayer.head()
```

Out[2]:

	player_id	long_scorecard_name	Salary	team	matches	wins	team_runs_for	team_runs_against	matches_keeper	byes_conceded	...	bowling_d
0	8931	AT Rayudu	343750.0	Chennai Super Kings	16	11	2809	2750	0	0	0	...
1	254771	D Shorey	31250.0	Chennai Super Kings	1	1	128	127	0	0	0	...
2	44613	DJ Bravo	1000000.0	Chennai Super Kings	16	11	2809	2750	0	0	0	...
3	214425	DJ Willey	NaN	Chennai Super Kings	3	2	484	483	0	0	0	...
4	258155	DL Chahar	125000.0	Chennai Super Kings	12	9	2117	2068	0	0	0	...

5 rows × 35 columns

◀ ▶

Data Exploration and Preparation

```
In [3]: IPLPlayer.shape
```

Out[3]: (149, 35)

Missing Values

```
In [4]: IPLPlayer.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 149 entries, 0 to 148
Data columns (total 35 columns):
player_id           149 non-null int64
long_scorecard_name 149 non-null object
Salary              141 non-null float64
team                149 non-null object
matches             149 non-null int64
wins                149 non-null int64
team_runs_for       149 non-null int64
team_runs_against   149 non-null int64
matches_keeper      149 non-null int64
byes_conceded       149 non-null int64
moms                149 non-null int64
innings             149 non-null int64
not_outs            149 non-null int64
runs                149 non-null int64
balls_faced         149 non-null int64
fours               149 non-null int64
sixes               149 non-null int64
matches_bowled      149 non-null int64
balls_bowled        149 non-null int64
wickets             149 non-null int64
runs_conceded       149 non-null int64
catches             149 non-null int64
stumpings           149 non-null int64
run_outs            149 non-null int64
batting_dot_balls   149 non-null int64
bowling_dot_balls   149 non-null int64
bowling_sixes       149 non-null int64
no_balls            149 non-null int64
balls_bowled_1_to_6 149 non-null int64
runs_conceded_1_to_6 149 non-null int64
balls_bowled_7_to_14 149 non-null int64
runs_conceded_7_to_14 149 non-null int64
balls_bowled_15_to_20 149 non-null int64
runs_conceded_15_to_20 149 non-null int64
event_winner         149 non-null int64
dtypes: float64(1), int64(32), object(2)
memory usage: 40.8+ KB
```

There are missing values in the salary variable. We will drop observations with missing values.

```
In [5]: IPLPlayer=IPLPlayer.dropna()
IPLPlayer.shape
```

```
Out[5]: (141, 35)
```

Create useful variables

Create dummy variables to indicate the role of the players.

- Create a variable to indicate whether a player had played as a batsman.

The variable "innings" indicates how many innings a player had batted in.

```
In [6]: IPLPlayer['batsman']=np.where(IPLPlayer['innings']> 0, 1, 0)
IPLPlayer['batsman'].describe()
```

```
Out[6]: count    141.000000
mean     0.943262
std      0.232165
min      0.000000
25%     1.000000
50%     1.000000
75%     1.000000
max     1.000000
Name: batsman, dtype: float64
```

- Create a variable to indicate bowler.

```
In [7]: IPLPlayer['bowler']=np.where(IPLPlayer['matches_bowled']> 0, 1, 0)
IPLPlayer['bowler'].describe()
```

```
Out[7]: count    141.000000
mean     0.631206
std      0.484198
min      0.000000
25%     0.000000
50%     1.000000
75%     1.000000
max     1.000000
Name: bowler, dtype: float64
```

The last type of player that is not captured by either batsman or bowler is wicket keeper. In the dataset, the variable "matches_keeper" indicates the number of matches that a player is a wicket keeper.

Performance Measures

1. batting average = runs / the numbers of outs
2. batting strike rate = (runs * 100) / balls faced
3. bowling average = runs conceded / wicket taken
4. bowling strike rate = number of balls bowled / wicket taken

Notice that if a batsman has scored runs but not been dismissed, his batting average is technically infinite. Similarly, if a player did not face any ball, his batting strike would be infinite and if a player did not lose any wicket, his bowling average or bowling strike would be infinite.

We will not be able to run a regression when our variables have some infinite values.

There are two alternatives we will consider to deal with this issue.

1. Add 1 to the number of outs, balls faced, andn wickets taken in calculating the above variables.
2. Instead of creating the above measures, we can simply include total runs, total number of outs, and balls faced to measure a batsman's performance, and include runs conceded, number of balls bowled, and wickets taken to measure a bowler's performance.

```
In [8]: IPLPlayer['outs']=np.where(IPLPlayer['batsman']==1, IPLPlayer['innings']-IPLPlayer['not_outs'], 0)
IPLPlayer['outs'].describe()
```

```
Out[8]: count    141.000000
mean     5.000000
std      4.605897
min      0.000000
25%     1.000000
50%     4.000000
75%     9.000000
max     16.000000
Name: outs, dtype: float64
```

Create batting average, batting stke rate, bowling average, and bowling strike rate variables. Add 1 to the number of outs, balls faced, andn wickets taken in calculating these variables.

```
In [9]: IPLPlayer['batting_average']=IPLPlayer['runs']/(IPLPlayer['outs']+1)
IPLPlayer['batting_strike']=IPLPlayer['runs']/((IPLPlayer['balls_faced']+1))*100
IPLPlayer['bowling_average']=IPLPlayer['runs_conceded']/(IPLPlayer['wickets']+1)
IPLPlayer['bowling_strike']=IPLPlayer['balls_bowled']/(IPLPlayer['wickets']+1)
```

```
In [10]: IPLPlayer['batting_average'].describe()
```

```
Out[10]: count    141.000000
mean     15.093066
std      13.761819
min      0.000000
25%     4.000000
50%     12.500000
75%     23.000000
max     65.000000
Name: batting_average, dtype: float64
```

```
In [11]: IPLPlayer['batting_strike'].describe()
```

```
Out[11]: count    141.000000
mean     104.164456
std      53.873378
min      0.000000
25%     73.913043
50%     118.446682
75%     139.669421
max     250.000000
Name: batting_strike, dtype: float64
```

```
In [12]: IPLPlayer['bowling_average'].describe()
```

```
Out[12]: count    141.000000
mean     17.493864
std      16.108488
min      0.000000
25%     0.000000
```

```

50%      20.052632
75%      27.466667
max      72.000000
Name: bowling_average, dtype: float64

```

```
In [13]: IPLPlayer['bowling_strike'].describe()

Out[13]:
count    141.000000
mean     11.478621
std      10.295591
min      0.000000
25%     0.000000
50%     12.500000
75%     19.600000
max     42.000000
Name: bowling_strike, dtype: float64

```

Regression Analyses

First let's run a regression of the salary on the type of player, batsman, bowler, and all-rounder.

```
In [14]: reg_IPL1=sm.ols(formula = 'Salary ~ batsman+ bowler+ batsman*bowler', data= IPLPlayer, missing="drop").fit()
print(reg_IPL1.summary())

OLS Regression Results
=====
Dep. Variable: Salary R-squared: 0.060
Model: OLS Adj. R-squared: 0.046
Method: Least Squares F-statistic: 4.379
Date: Mon, 26 Jul 2021 Prob (F-statistic): 0.0143
Time: 00:22:00 Log-Likelihood: -2069.2
No. Observations: 141 AIC: 4144.
Df Residuals: 138 BIC: 4153.
Df Model: 2
Covariance Type: nonrobust
=====

coef std err t P>|t| [0.025 0.975]
-----
Intercept 2.859e+05 1.11e+05 2.577 0.011 6.65e+04 5.05e+05
batsman 4.61e+05 1.11e+05 4.156 0.000 2.42e+05 6.8e+05
bowler -1.668e+05 1.11e+05 -1.584 0.135 -3.86e+05 5.25e+04
batsman:bowler 8340.1549 1.2e+05 0.070 0.945 -2.29e+05 2.45e+05
Omnibus: 29.772 Durbin-Watson: 2.052
=====
```

Next we will first focus on performance of batsman.

We will first simply use the total number of runs, number of not outs, and number of balls faced to measure players' performance.

```
In [15]: reg_IPL2=sm.ols(formula = 'Salary ~ runs', data= IPLPlayer).fit()
print(reg_IPL2.summary())

OLS Regression Results
=====
Dep. Variable: Salary R-squared: 0.267
Model: OLS Adj. R-squared: 0.261
Method: Least Squares F-statistic: 50.57
Date: Mon, 26 Jul 2021 Prob (F-statistic): 5.54e-11
Time: 00:22:01 Log-Likelihood: -2051.7
No. Observations: 141 AIC: 4107.
Df Residuals: 139 BIC: 4113.
Df Model: 1
Covariance Type: nonrobust
=====

coef std err t P>|t| [0.025 0.975]
-----
Intercept 3.878e+05 5.39e+04 7.198 0.000 2.81e+05 4.94e+05
runs 1737.9337 244.4800 7.111 0.000 1254.711 2221.157
Omnibus: 19.056 Durbin-Watson: 2.116
Prob(Omnibus): 0.000 Jarque-Bera (JB): 22.155
Skew: 0.913 Prob(JB): 1.55e-05
Kurtosis: 3.660 Cond. No. 277.
=====

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
```

```
In [16]: reg_IPL3=sm.ols(formula = 'Salary ~ runs+not_outs', data= IPLPlayer).fit()
print(reg_IPL3.summary())

OLS Regression Results
=====
Dep. Variable: Salary R-squared: 0.318
Model: OLS Adj. R-squared: 0.308
Method: Least Squares F-statistic: 32.15
Date: Mon, 26 Jul 2021 Prob (F-statistic): 3.45e-12
Time: 00:22:01 Log-Likelihood: -2046.6
No. Observations: 141 AIC: 4099.
Df Residuals: 138 BIC: 4108.
Df Model: 2
Covariance Type: nonrobust
=====

coef std err t P>|t| [0.025 0.975]
-----
Intercept 2.88e+05 6.07e+04 4.747 0.000 1.68e+05 4.08e+05
runs 1491.1582 248.725 5.995 0.000 999.393 1982.963
not_outs 8.955e+04 2.79e+04 3.215 0.002 3.45e+04 1.45e+05
Omnibus: 16.162 Durbin-Watson: 2.092
Prob(Omnibus): 0.000 Jarque-Bera (JB): 18.001
Skew: 0.886 Prob(JB): 0.000123
Kurtosis: 3.682 Cond. No. 333.
=====

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
```

```
In [17]: reg_IPL4=sm.ols(formula = 'Salary ~ runs+not_outs+balls_faced', data= IPLPlayer).fit()
print(reg_IPL4.summary())

OLS Regression Results
=====
Dep. Variable: Salary R-squared: 0.321
Model: OLS Adj. R-squared: 0.306
Method: Least Squares F-statistic: 21.60
Date: Mon, 26 Jul 2021 Prob (F-statistic): 1.62e-11
Time: 00:22:01 Log-Likelihood: -2046.3
No. Observations: 141 AIC: 4101.
```

```

Df Residuals: 137 BIC: 4112.
Df Model: 3
Covariance Type: nonrobust
=====
      coef    std err      t   P>|t|      [0.025    0.975]
-----
Intercept 3.013e+05 6.29e+04 4.791  0.000  1.77e+05  4.26e+05
runs     2871.9166 1711.997 1.678  0.096 -513.440  6257.273
not_outs 8.945e+04 2.79e+04 3.207  0.002  3.43e+04  1.45e+05
balls_faced -2044.5867 2508.113 -0.815  0.416 -7004.207  2915.033
=====
Omnibus: 16.922 Durbin-Watson: 2.095
Prob(Omnibus): 0.000 Jarque-Bera (JB): 19.142
Skew: 0.819 Prob(JB): 6.97e-05
Kurtosis: 3.758 Cond. No. 418.
=====
```

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In the next regressions, we will use the modified batting average and batting strike variables to measure player performance.

```

In [18]: reg_IPL5=sm.ols(formula = 'Salary ~ batting_average', data=IPLPlayer).fit()
print(reg_IPL5.summary())

OLS Regression Results
=====
Dep. Variable: Salary R-squared: 0.233
Model: OLS Adj. R-squared: 0.227
Method: Least Squares F-statistic: 42.13
Date: Mon, 26 Jul 2021 Prob (F-statistic): 1.40e-09
Time: 00:22:01 Log-Likelihood: -2054.9
No. Observations: 141 AIC: 4114.
Df Residuals: 139 BIC: 4120.
Df Model: 1
Covariance Type: nonrobust
=====
      coef    std err      t   P>|t|      [0.025    0.975]
-----
Intercept 3.072e+05 6.52e+04 4.716  0.000  1.78e+05  4.36e+05
batting_average 2.074e+04 3194.873 6.491  0.000  1.44e+04  2.71e+04
=====
Omnibus: 18.526 Durbin-Watson: 2.083
Prob(Omnibus): 0.000 Jarque-Bera (JB): 21.526
Skew: 0.929 Prob(JB): 2.12e-05
Kurtosis: 3.457 Cond. No. 30.4
=====
```

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

In [19]: reg_IPL6=sm.ols(formula = 'Salary ~ batting_average+batting_strike', data=IPLPlayer).fit()
print(reg_IPL6.summary())

OLS Regression Results
=====
Dep. Variable: Salary R-squared: 0.234
Model: OLS Adj. R-squared: 0.223
Method: Least Squares F-statistic: 21.12
Date: Mon, 26 Jul 2021 Prob (F-statistic): 9.96e-09
Time: 00:22:02 Log-Likelihood: -2054.7
No. Observations: 141 AIC: 4115.
Df Residuals: 138 BIC: 4124.
Df Model: 2
Covariance Type: nonrobust
=====
      coef    std err      t   P>|t|      [0.025    0.975]
-----
Intercept 2.668e+05 9.69e+04 2.754  0.007  7.52e+04  4.58e+05
batting_average 1.903e+04 4409.440 4.315  0.000  1.03e+04  2.77e+04
batting_strike 635.3041 1126.380 0.564  0.574 -1591.892  2862.500
=====
Omnibus: 17.736 Durbin-Watson: 2.061
Prob(Omnibus): 0.000 Jarque-Bera (JB): 20.405
Skew: 0.988 Prob(JB): 3.71e-05
Kurtosis: 3.416 Cond. No. 262.
=====
```

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

We will now turn to bowlers' performance.

Again, we will first use number of runs conceded, number of balls bowled, and number of wickets taken to measure bowlers' performance.

```

In [20]: reg_IPL7=sm.ols(formula = 'Salary ~ runs_conceded', data=IPLPlayer).fit()
print(reg_IPL7.summary())

OLS Regression Results
=====
Dep. Variable: Salary R-squared: 0.023
Model: OLS Adj. R-squared: 0.015
Method: Least Squares F-statistic: 3.200
Date: Mon, 26 Jul 2021 Prob (F-statistic): 0.0758
Time: 00:22:02 Log-Likelihood: -2072.0
No. Observations: 141 AIC: 4148.
Df Residuals: 139 BIC: 4154.
Df Model: 1
Covariance Type: nonrobust
=====
      coef    std err      t   P>|t|      [0.025    0.975]
-----
Intercept 5.438e+05 6.53e+04 8.322  0.000  4.15e+05  6.73e+05
runs_conceded 569.3989 318.271 1.789  0.076 -59.968  1198.587
=====
Omnibus: 39.245 Durbin-Watson: 2.073
Prob(Omnibus): 0.000 Jarque-Bera (JB): 62.394
Skew: 1.414 Prob(JB): 2.83e-14
Kurtosis: 4.621 Cond. No. 271.
=====
```

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

In [21]: reg_IPL8=sm.ols(formula = 'Salary ~ runs_conceded+balls_bowled', data=IPLPlayer).fit()
print(reg_IPL8.summary())

OLS Regression Results
=====
Dep. Variable: Salary R-squared: 0.042
```

```

Model: OLS Adj. R-squared: 0.028
Method: Least Squares F-statistic: 3.026
Date: Mon, 26 Jul 2021 Prob (F-statistic): 0.0518
Time: 00:22:02 Log-Likelihood: -2070.5
No. Observations: 141 AIC: 4147.
Df Residuals: 138 BIC: 4156.
Df Model: 2
Covariance Type: nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	5.565e+05	6.54e+04	8.514	0.000	4.27e+05	6.86e+05
runs conceded	-2749.4793	2004.845	-1.371	0.172	-6713.666	1214.707
balls_bowled	4574.8787	2729.042	1.676	0.096	-821.265	9971.023

```

Omnibus: 41.062 Durbin-Watson: 2.075
Prob(Omnibus): 0.000 Jarque-Bera (JB): 67.476
Skew: 1.449 Prob(JB): 2.23e-15
Kurtosis: 4.756 Cond. No. 336.
=====
```

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [22]: reg_IPL9=sm.ols(formula = 'Salary ~ runs_conceded+balls_bowled+wickets', data= IPLPlayer).fit()
print(reg_IPL9.summary())
```

```

OLS Regression Results
=====
Dep. Variable: Salary R-squared: 0.049
Model: OLS Adj. R-squared: 0.028
Method: Least Squares F-statistic: 2.329
Date: Mon, 26 Jul 2021 Prob (F-statistic): 0.0772
Time: 00:22:02 Log-Likelihood: -2070.1
No. Observations: 141 AIC: 4148.
Df Residuals: 137 BIC: 4160.
Df Model: 3
Covariance Type: nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	5.543e+05	6.54e+04	8.472	0.000	4.25e+05	6.84e+05
runs conceded	-3049.2063	2029.065	-1.503	0.135	-7061.542	963.129
balls_bowled	6343.2378	3284.524	1.931	0.056	-151.683	1.28e+04
wickets	-2.737e+04	2.83e+04	-0.968	0.335	-8.33e+04	2.85e+04

```

Omnibus: 41.167 Durbin-Watson: 2.050
Prob(Omnibus): 0.000 Jarque-Bera (JB): 67.717
Skew: 1.453 Prob(JB): 1.97e-15
Kurtosis: 4.757 Cond. No. 337.
=====
```

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In the next regression, we will use the modified bowling average and bowling strike variables to measure player performance.

```
In [23]: reg_IPL10=sm.ols(formula = 'Salary ~ bowling_average+bowling_strike', data= IPLPlayer).fit()
print(reg_IPL10.summary())
```

```

OLS Regression Results
=====
Dep. Variable: Salary R-squared: 0.054
Model: OLS Adj. R-squared: 0.040
Method: Least Squares F-statistic: 3.912
Date: Mon, 26 Jul 2021 Prob (F-statistic): 0.0223
Time: 00:22:03 Log-Likelihood: -2069.7
No. Observations: 141 AIC: 4145.
Df Residuals: 138 BIC: 4154.
Df Model: 2
Covariance Type: nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	6.535e+05	7.33e+04	8.918	0.000	5.09e+05	7.98e+05
bowling_average	-3.415e+04	1.25e+04	-2.740	0.007	-5.88e+04	-9587.025
bowling_strike	4.914e+04	1.95e+04	2.520	0.013	1.06e+04	8.77e+04

```

Omnibus: 33.704 Durbin-Watson: 2.040
Prob(Omnibus): 0.000 Jarque-Bera (JB): 48.909
Skew: 1.293 Prob(JB): 2.40e-11
Kurtosis: 4.280 Cond. No. 42.4
=====
```

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Lastly, we will incorporate performance measures of both batsman and bowler in the same regression.

We will first use the original variables, total number of runs, number of not outs, number of balls faced, number of runs conceded, number of balls bowled, and number of wickets in the regression.

```
In [24]: reg_IPL11=sm.ols(formula = 'Salary ~ runs+not_outs+balls_faced+runs_conceded+balls_bowled+wickets', data= IPLPlayer).fit()
print(reg_IPL11.summary())
```

```

OLS Regression Results
=====
Dep. Variable: Salary R-squared: 0.408
Model: OLS Adj. R-squared: 0.382
Method: Least Squares F-statistic: 15.41
Date: Mon, 26 Jul 2021 Prob (F-statistic): 2.20e-13
Time: 00:22:03 Log-Likelihood: -2036.6
No. Observations: 141 AIC: 4087.
Df Residuals: 134 BIC: 4108.
Df Model: 6
Covariance Type: nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.458e+05	7.32e+04	1.993	0.048	1123.895	2.91e+05
runs	2089.9774	1629.015	1.283	0.202	-1131.930	5311.885
not_outs	5.95e+04	2.82e+04	2.110	0.037	3731.096	1.15e+05
balls_faced	-354.0166	2412.458	-0.147	0.884	-5125.422	4417.389
runs_conceded	-1737.5721	1672.019	-1.039	0.301	-5044.535	1569.391
balls_bowled	5030.2797	2648.309	1.899	0.060	-207.614	1.03e+04
wickets	-2.231e+04	2.26e+04	-0.989	0.325	-6.69e+04	2.23e+04

```

Omnibus: 17.606 Durbin-Watson: 2.049
Prob(Omnibus): 0.000 Jarque-Bera (JB): 23.546
Skew: 0.709 Prob(JB): 7.71e-06
=====
```

```
Kurtosis: 4.412 Cond. No. 538.
```

```
=====  
Warnings:  
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
```

We will also use the modified batting average, batting strike, bowling average, and bowling strike variables to measure the player performance.

```
In [25]: reg_IPL12=sm.ols(formula = 'Salary ~ modified_batting_average+modified_batting_strike+modified_bowling_average+modified_bowling_strike', data= IPLPlayer).fit()  
print(reg_IPL12.summary())
```

```
OLS Regression Results  
=====  
Dep. Variable: Salary R-squared: 0.308  
Model: OLS Adj. R-squared: 0.288  
Method: Least Squares F-statistic: 15.16  
Date: Mon, 26 Jul 2021 Prob (F-statistic): 2.85e-10  
Time: 00:22:04 Log-Likelihood: -2047.6  
No. Observations: 141 AIC: 4105.  
Df Residuals: 136 BIC: 4120.  
Df Model: 4  
Covariance Type: nonrobust  
=====  
coef std err t P>|t| [0.025 0.975]  
-----  
Intercept 1.37e+05 1.14e+05 1.204 0.231 -8.81e+04 3.62e+05  
modified_batting_average 2.418e+04 4617.989 5.237 0.000 1.51e+04 3.33e+04  
modified_batting_strike -61.2223 1094.543 -0.056 0.955 -2225.747 2103.303  
modified_bowling_average -3.186e+04 1.08e+04 -2.951 0.004 -5.32e+04 -1.05e+04  
modified_bowling_strike 5.941e+04 1.7e+04 3.492 0.001 2.58e+04 9.31e+04  
=====  
Omnibus: 13.575 Durbin-Watson: 2.103  
Prob(Omnibus): 0.001 Jarque-Bera (JB): 14.473  
Skew: 0.732 Prob(JB): 0.000720  
Kurtosis: 3.565 Cond. No. 324.  
=====
```

```
Warnings:  
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
```

Self Test

- Run a regression of salary as a function of the interaction of batsmen and runs and the interaction of bowler and wickets taken.
- Interpret your regression results.

```
In [ ]:
```